

### 0.0.1 Question 1a

In this project we are essentially trying to answer the question. “How much is a house worth?” - Who might be interested in an answer to this question? **Please list at least three different parties (people or organizations) and then describe whether each one has an interest in seeing the housing price to be either high or low.**

Parties that may be interested in trying to answer this question are: real estate agents, homeowners, and prospective home buyers. Real estate agents would be interested in seeing and keeping track of housing prices in order to determine when to sell to clients. When housing prices are higher, agents are able to sell for a higher profit which allows them to receive a higher commission. As for homeowners, if they want to sell their existing home, they would seek to sell when housing prices are higher to earn a profit. On the contrary, home buyers would look to purchase when housing prices are low since they're more affordable and likely have lower interest rates as well.



### 0.0.2 Question 1b

- 1bi). Which of the following scenarios strike you as unfair and why? You can choose more than one. There is no single right answer, but you must explain your reasoning.
- 1bii). Would you consider some of these scenarios more (or less) fair than others? Why?

Scenario A: A homeowner whose home is assessed at a higher price than it would sell for.

Scenario B: A homeowner whose home is assessed at a lower price than it would sell for.

Scenario C: An assessment process that systematically overvalues inexpensive properties and undervalues expensive properties.

Scenario D: An assessment process that systematically undervalues inexpensive properties and overvalues expensive properties.

Write your full answers to both parts in the cell below:

All of these scenarios are seemingly unfair due to the fact that they either over or undervalue some aspect of each property. Scenario C is the most unfair due to its systematic bias against lower-value property owners since it places a greater burden on those who are less able to bear it, while Scenario A is also concerning for the financial burden it has on people. Scenario B and Scenario D may seem less unfair, but they're still concerning when it comes to equity and systemic impacts.



### 0.0.3 Question 1d

- 1di). What were the central problems with the earlier property tax system in Cook County as reported by the Chicago Tribune?
- 1dii). What were the primary causes of these problems? (Note: in addition to reading the paragraph above you will need to **read the [Project Case Study.pdf](#) explaining the context and history of this dataset before answering this question**).

1di.) The central problems involved biased property tax assessments, which undervalued high-priced homes and overvaluing low-priced homes. This led to wealthier homeowners paying less than their fair share of property taxes, while working-class homeowners paid more. Similarly, there were racial disparities where white homeowners benefited from the undervaluations, while working-class homeowners (people of color) were unfairly taxed at higher rates. Assessments were also consistently “off the mark” for years, demonstrating a lack of precision and fairness in the valuation process.

1dii.) Primary causes of these problems seemed to stem from the fact that the old system relied on outdated or inconsistent methods that failed to accurately estimate property values. Errors varied across different groups, regions, or property types, leading to unequal treatment. Higher-value properties were also often under-assessed compared to lower-value properties, which resulted in disproportionate tax burdens on lower income homeowners. Additionally, the Board of Review appeals process created opportunities for favoritism and unfair influence. Wealthier property owners could afford to appeal their assessments, which often led to reduced tax burdens, further skewing equity. Existing data was incomplete, outdated, or biased. For instance, missing data was more prevalent in lower-income neighborhoods, and renovations or improvements were inconsistently reported.



#### 0.0.4 Question 2b:

Since we're trying to predict Sale Price, next we'll look for missing or unusual outliers in that field.

Examine the Sale Price column in the `training_val_data` DataFrame and answer the following questions:

- 2bi). Does the Sale Price data have any missing, N/A, negative or 0 values for the data? If so, propose a way to handle this.
- 2bii). Does the Sale Price data have any unusually large outlier values? If so, propose a cutoff to use for throwing out large outliers, and justify your reasoning).
- 2biii). Does the Sale Price data have any unusually small outlier values? If so, propose a cutoff to use for throwing out small outliers, and justify your reasoning.

Below are three cells. The first is a Markdown cell for you to write up your responses to all 3 parts above. The second two are code cells that are available for you to write code to explore the outliers and/or visualize the Sale Price data.

#### 0.0.5 Question 2b i, ii, iii answer cell: *Type your responses to all three parts in this cell...*

2bi.) Based on my calculations, it appears that Sale Price does not have any missing or negative/0 values within the data.

2bii.) Sale Price does have some unusually large outlier values, with prices exceeding \$712200. In order to combat this, the rows containing these high sale price can be removed to not interfere with the rest of the data, or we could limit the data to only examine properties that fall under this sale threshold.

2biii.) Similarly, it appears there are some unusually low prices as \$1 was the lower limit threshold, but no house sells for only \$1. So similarly, we could just remove all the rows that contain absurdly low sale prices so it doesn't conflict with the general data set.

```
In [12]: missing_values = tr_val_data['Sale Price'].isna().sum()

negative_or_zero_values = (tr_val_data['Sale Price'] <= 0).sum()

print("Missing or N/A values:", missing_values)
print("Negative or 0 values:", negative_or_zero_values)
# your code exploring Sale Price above this line
```

```
Missing or N/A values: 0
Negative or 0 values: 0
```

```

In [13]: Q1 = tr_val_data['Sale Price'].quantile(0.25)
        Q3 = tr_val_data['Sale Price'].quantile(0.75)
        IQR = Q3 - Q1
        upper_limit = Q3 + 1.5 * IQR
        large_outliers = tr_val_data[tr_val_data['Sale Price'] > upper_limit]

        print("Number of large outliers:", large_outliers.shape[0])
        print("Upper limit for outliers:", upper_limit)

        # Calculate lower quartile (Q1) and other percentiles
        Q1 = tr_val_data['Sale Price'].quantile(0.25)
        lower_limit = tr_val_data['Sale Price'].quantile(0.05)

        small_outliers = tr_val_data[tr_val_data['Sale Price'] < lower_limit]

        print("Number of small outliers:", len(small_outliers))
        print("Lower limit for small outliers:", lower_limit)

        # optional extra cell for exploring code

```

```

Number of large outliers: 12229
Upper limit for outliers: 712200.0
Number of small outliers: 0
Lower limit for small outliers: 1.0

```



### 0.0.6 Question 6a: Choose an additional feature

It's your turn to choose another feature to add to the model. Choose one new **quantitative** (not qualitative) feature and create Model 3 incorporating this feature (along with the features we've already chosen in Model 2). Try to choose a feature that will have a large impact on reducing the RMSE and/or will improve your residual plots. This can be a raw feature available in the dataset, or a transformation of one of the features in the dataset, or a new feature that you create from the dataset (see Project 1 for ideas).

Note: There is not one single right answer as to which feature to add, however **to receive credit on this question you should make sure the feature decreases the Cross Validation RMSE compared to Model 2 (i.e. we want to improve the model, not make it worse!)**

In the cell below, explain what additional feature you have chosen and why. Justify your reasoning. There are optional code cells provided below for you to use when exploring the dataset to determine which feature to add.

This problem will be graded based on your reasoning and explanation of the feature you choose, and then on your implementation of incorporating the feature.

**NOTE** Please don't add additional coding cells below or the Autograder will have issues. You do not need to use all the coding cells provided.

### 0.0.7 Question 6a Answer Cell:

In this cell, explain what feature you chose to add and why. Then give the equation for your new model (use the LaTeX from Model 2 from above and then add an additional term).

$$\text{Log Sale Price} = \theta_1(\text{Log Building Square Feet}) + \theta_2(\text{Shingle/Asphalt})$$

$$+ \theta_3(\text{Tar\&Gravel}) + \theta_4(\text{Tile}) + \theta_5(\text{Shake}) + \theta_6(\text{Other}) + \theta_7(\text{Slate}) + \theta_8(\text{Estimate (Building)})$$

I was originally going to incorporate the Age Decade column as the age of the house usually plays a decent role in determining the property's overall value. Depending on the property, age could either increase or decrease a property's value due to depreciation or vintage styling choices. Similarly, older houses have more wear and tear which may also decrease its overall value, so it seemed like a good choice to help improve the model. However, once I got to calculating the RMSE, the value was higher than before, so I chose the next qualitative column I saw: Neighborhood Code. However, that also was not optimal in improving my RMSE value, which led me to use Building (Estimate). Building (Estimate) seems to be the estimated sale price for the property, so it sounds like it would also be a good feature to add to the model since we're dealing with sale prices.

```
In [53]: tr_val_data["Age Decade"].unique()
```

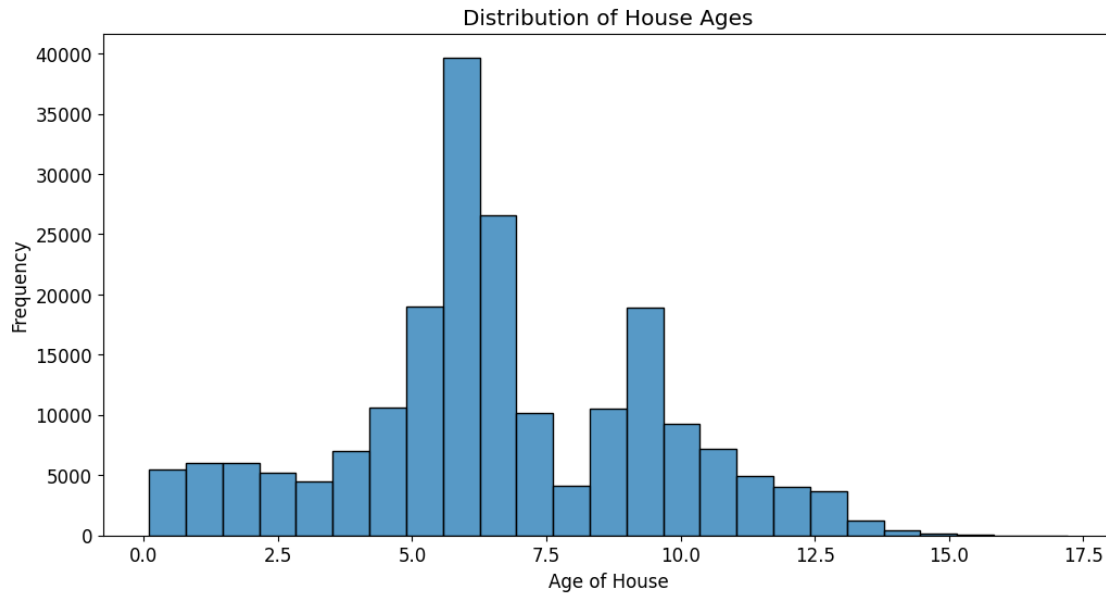
*# Show work in this cell exploring data to determine which feature to add*

```
Out[53]: array([13.2,  9.6, 11.2,  6.3,  5.8, 10.9,  1.7, 10. ,  4.8,  7.4,  3.4,
                1.3, 12.2,  1.6,  5.9,  9.4,  8.7,  4.1,  6.5,  6.9,  0.1,  6.4,
                9.5,  2.7,  9.2,  6.2,  7.3,  6.7, 10.7,  9.3,  5.4,  3.3,  4. ,
                0.7,  9.1,  4.2,  1. ,  3.8,  2.8,  2.5, 10.2,  4.9,  5.7,  6.1,
                3. ,  3.9,  6.6, 10.1,  7.1,  4.7,  0.9,  5.1,  6. ,  8.8, 10.5,
                7.5,  8.9,  7.7,  4.4,  4.6, 11.5,  9. ,  6.8, 11.7,  4.3, 10.6,
                5.5, 11. ,  8.5,  3.7,  7.6, 13.4,  8. , 12.9,  4.5,  5. ,  1.9,
                8.4,  5.3,  3.5,  0.8,  0.4, 12.3, 14. ,  5.6,  7.8,  3.6, 12.1,
                11.8,  8.6,  3.1, 12.7, 12.8,  1.5,  2.9,  5.2,  1.8,  8.3,  8.2,
                10.4,  9.8,  9.7, 11.4, 11.3,  9.9,  2.4,  7. , 12. , 11.6, 11.1,
                12.5,  3.2, 12.4, 13.6,  7.2,  2.1, 10.3,  2.2,  2. , 10.8, 13.3,
                0.6, 15. , 13. ,  0.5,  1.2,  7.9, 13.7,  1.4,  2.6, 14.3,  1.1,
                11.9,  8.1, 12.6, 14.7, 14.4, 15.1, 13.5,  2.3, 14.2, 13.1, 13.9,
                13.8, 15.5, 14.8, 14.1, 14.5, 15.9,  0.3, 17.2, 14.6, 15.4, 15.6,
                14.9, 15.2, 15.3,  0.2, 16. , 16.5, 16.1, 16.3, 16.9, 16.4, 15.8,
                15.7])
```

```
In [54]: display(tr_val_data['Age Decade'].describe())
plt.figure(figsize=(12, 6))
sns.histplot(tr_val_data['Age Decade'], bins = 25)
plt.title("Distribution of House Ages")
plt.xlabel("Age of House")
plt.ylabel("Frequency")
plt.show()
```

*# Optional code cell for additional work exploring data/ explaining which feature you chose.*

```
count    204792.000000
mean         6.598121
std         2.900149
min          0.100000
25%         5.100000
50%         6.200000
75%         8.900000
max        17.200000
Name: Age Decade, dtype: float64
```



```
In [55]: tr_val_data['Neighborhood Code'].unique()
```

*# Optional code cell for additional work exploring data/ explaining which feature you chose.*

```
Out[55]: array([ 50, 120, 210, 220, 380, 181,  52,  70,  33,  20,  21,  30, 100,
                40,  61,  24, 282,  11, 112, 130, 240,  60,  10,  41,  92, 330,
                31,  32,  93,  51, 520,  84,  85,  88, 180, 110, 194, 160,  71,
                171,  22, 141, 113,  80,  13, 440, 212, 200,  45, 250, 190, 274,
                150,  35, 422, 162, 142, 260,  81, 314,  82,  42, 361, 430,  74,
                410, 103, 420, 321, 102,  12,  23, 132, 122, 101, 350, 151, 223,
                115,  90, 323,  91,  63, 270, 281,  15, 280, 560,  14, 170, 211,
                111, 191, 271,  55, 251, 104, 221, 241,  53, 371,  72, 461,  75,
                342, 345, 175, 320,  43,  62, 423, 140, 230, 362,  94,  34, 193,
                310, 121,  54, 390,  46, 344, 114,  19,  44,  65, 431, 166, 131,
                47, 312, 275, 161,  39, 257,  83, 255, 183, 192,  73,  36, 222,
                316,  96, 164, 315, 580,  64, 300,  26, 224,  87, 226, 400, 152,
                56,  25, 293, 402, 182, 174, 340, 201, 133, 116, 185,  27, 600,
                432, 165,  48,  37,  67, 360, 262, 232, 134,  99, 117, 227, 290,
                163,  18, 109,  95,  38, 143,  86, 145, 463, 341, 106])
```

```
In [56]: tr_val_data['Estimate (Building)'].unique()
```

*# Optional code cell for additional work exploring data/ explaining which feature you chose.*

```
Out[56]: array([139500, 177500,  63470, ..., 635930, 537990, 637720])
```



### 0.0.8 Question 6b: Create Model 3

In the cells below fill in the code to create and analyze Model 3 (follow the Modeling steps outlined above).

PLEASE DO NOT ADD ANY ADDITIONAL CELLS IN THIS PROBLEM OR IT MIGHT MAKE THE AUTOGRADER FAIL

```
In [57]: # Modeling Step 1: Process the Data

# Hint: You can either use your implementation of the One Hot Encoding Function
# from Project Part 1, or use the staff's implementation

from feature_func import *

...
# Optional: Define any helper functions you need for one-hot encoding above this line

def process_data_m3(df):

    data = df.copy()

    data = data[data["Pure Market Filter"]==1]

    # Create Log Sale Price column
    data["Log Sale Price"] = np.log(data["Sale Price"])

    # Create Log Building Square Feet column
    data["Log Building Square Feet"] = np.log(data["Building Square Feet"])

    # Create Log Estimate (Building) column
    data["Log Estimate (Building)"] = np.log(data["Estimate (Building)"] + 1)

    # Change Roof Material to names
    data = substitute_roof_material(data)

    # one-hot encode the roof material
    data = ohe_roof_material(data)

    # select only relevant columns
    relevant_columns = ['Log Estimate (Building)', 'Log Building Square Feet', 'Log Sale Price']
    data = data[relevant_columns]

    return data

# Process the data for Model 3 (using the same tr and val datasets we created in Question 3)
processed_train_m3 = process_data_m3(tr)
```

```

processed_val_m3 = process_data_m3(val)

# Create X (Dataframe) and y (series) to use to train the model
X_train_m3 = processed_train_m3.drop(columns = ['Log Sale Price'])
y_train_m3 = processed_train_m3['Log Sale Price']

X_valid_m3 = processed_val_m3.drop(columns = ['Log Sale Price'])
y_valid_m3 = processed_val_m3['Log Sale Price']

# Take a look at the result
display(X_train_m3.head())
display(y_train_m3.head())

display(X_valid_m3.head())
display(y_valid_m3.head())

```

	Log Estimate (Building)	Log Building Square Feet \
21302	11.248061	6.871091
19451	11.510031	7.576610
32018	12.096431	6.891626
144262	12.042265	7.186901
197227	12.612640	7.576610

	Roof Material_Other	Roof Material_Shake \
21302	0.0	0.0
19451	0.0	0.0
32018	0.0	0.0
144262	0.0	0.0
197227	0.0	0.0

	Roof Material_Shingle/Asphalt	Roof Material_Slate \
21302	1.0	0.0
19451	1.0	0.0
32018	1.0	0.0
144262	1.0	0.0
197227	1.0	0.0

	Roof Material_Tar&Gravel	Roof Material_Tile
21302	0.0	0.0
19451	0.0	0.0
32018	0.0	0.0
144262	0.0	0.0
197227	0.0	0.0

21302	11.738466
19451	12.100712
32018	12.577291
144262	12.502467
197227	12.449019

Name: Log Sale Price, dtype: float64

	Log Estimate (Building)	Log Building Square Feet	\
17112	11.882217	7.487174	
189337	12.110118	7.688913	
141725	11.760418	6.985642	
9776	11.569599	6.846943	
81676	11.272776	7.096721	

	Roof Material_Other	Roof Material_Shake	\
17112	0.0	0.0	
189337	0.0	0.0	
141725	0.0	0.0	
9776	0.0	0.0	
81676	0.0	0.0	

	Roof Material_Shingle/Asphalt	Roof Material_Slate	\
17112	1.0	0.0	
189337	1.0	0.0	
141725	1.0	0.0	
9776	1.0	0.0	
81676	1.0	0.0	

	Roof Material_Tar&Gravel	Roof Material_Tile
17112	0.0	0.0
189337	0.0	0.0
141725	0.0	0.0
9776	0.0	0.0
81676	0.0	0.0

17112	12.100707
189337	12.460715
141725	12.165251
9776	11.982929
81676	11.719940

Name: Log Sale Price, dtype: float64

In [58]: # Modeling STEP 2: Create and Fit a Multiple Linear Regression Model

```
linear_model_m3 = lm.LinearRegression(fit_intercept=False)

linear_model_m3.fit(X_train_m3, y_train_m3)

# your code above this line to create and fit regression model for Model 3

y_predict_train_m3 = linear_model_m3.predict(X_train_m3)

y_predict_valid_m3 = linear_model_m3.predict(X_valid_m3)
```

In [59]: # MODELING STEP 3: Evaluate the RMSE for your model

```

# Training and validation errors for the model (in units dollars, not log(dollars))

training_error[2] = rmse(y_predict_train_m3, y_train_m3)
validation_error[2] = rmse(y_predict_valid_m3, y_valid_m3)

training_error[2] = rmse(np.exp(y_predict_train_m3), np.exp(y_train_m3))
validation_error[2] = rmse(np.exp(y_predict_valid_m3), np.exp(y_valid_m3))

(print("3rd Model \nTraining RMSE: $ {}".format(training_error[2], validation_error[2]))
)

```

3rd Model  
Training RMSE: \$ 232349.50469106025  
Validation RMSE: 234597.11998179482

In [60]: # MODELING STEP 4: Conduct 5-fold cross validation for model and output RMSE

```

linear_model_m3_cv = lm.LinearRegression(fit_intercept=False)

# Process the entire cleaned training_val dataset using the m3 pipeline
processed_full_m3 = process_data_m3(tr_val_data)

# Split the processed_full_m3 Dataset into X and y to use in models.
X_full_m3 = processed_full_m3.drop(columns="Log Sale Price")
y_full_m3 = processed_full_m3["Log Sale Price"]

# Run cross_validate_rmse function:
cv_error_m3 = cross_validate_rmse(linear_model_m3_cv, X_full_m3, y_full_m3)

# Save the cross validation error for model 3 in our list to compare different models:

cv_error[2] = cv_error_m3

print("3rd Model Cross Validation RMSE: {}".format(cv_error[2]))

```

3rd Model Cross Validation RMSE: 232477.29879016159

In [61]: # MODELING STEP 5: Add a name for your 3rd model describing the features  
#and run this cell to Plot bar graph all 3 models

```

model_names[2] = "M3: log(bsqft)+log(estbldg)+Roof"


```



```
fig = go.Figure([
go.Bar(x = model_names, y = training_error, name="Training RMSE"),
go.Bar(x = model_names, y = validation_error, name="Validation RMSE"),
go.Bar(x = model_names, y = cv_error, name="Cross Val RMSE")
])
```

```
fig.update_yaxes(range=[180000,260000], title="RMSE")
```

```
fig
```



```
In [62]: # MODELING STEP 5 cont'd: Plot 2 side-by-side residual plots
#(similar to Question 3, for validation data)
```

```
fig, ax = plt.subplots(1,2, figsize=(15, 5))
```

```
residuals = y_valid_m3 - y_predict_valid_m3
```

```
x_plt1 = y_predict_valid_m3
y_plt1 = residuals
```

```
x_plt2 = y_valid_m3
y_plt2 = residuals
```

```
ax[0].scatter(x_plt1, y_plt1, alpha=.25)
ax[0].axhline(0, c='black', linewidth=1)
ax[0].set_xlabel(r'Predicted Log(Sale Price)')
ax[0].set_ylabel(r'Residuals: Log(Sale Price) - Predicted Log(Sale Price)');
ax[0].set_title("Model 3 Val Data: Residuals vs. Predicted Log(Sale Price)")
```


```
ax[1].scatter(x_plt2, y_plt2, alpha=.25)
```

```

ax[1].axhline(0, c='black', linewidth=1)
ax[1].set_xlabel(r'Log(Sale Price)')
ax[1].set_ylabel(r'Residuals: Log(Sale Price) - Predicted Log(Sale Price)');
ax[1].set_title("Model 3 Val Data: Residuals vs. Log(Sale Price)")

```

Out[62]: Text(0.5, 1.0, 'Model 3 Val Data: Residuals vs. Log(Sale Price)')



In [63]: grader.check("q6b")

Out[63]: q6b results: All test cases passed!

### 0.0.9 Question 6c

- 6ci). Comment on your RMSE and residual plots from Model 3 compared to the first 2 models.
- 6cii). Are the residuals of your model still showing a trend that overestimates lower priced houses and underestimates higher priced houses? If so, how could you try to address this in the next round of modeling?
- 6ciii). If you had more time to improve your model, what would your next steps be?

6ci). The RMSE and residual plots from Model 3 seem to be better than the first two models. Since the distribution is more evenly distributed around 0, we can tell that the estimation of houses is more accurate to their actual values.

6cii). There is still a slight trend which we can try and accomodate for by adding more features in future models. For instance, we could try and incorporate other features that may be closely correlated with the sale price so as to avoid the over and underestimation.

6ciii). If I had more time to improve my model, I would go back and try to incorporate age as I feel that would be a good factor. I was originally going to incorporate that, but messed up somewhere so with more time, I would try and debug it to see if it would work.



---

## 0.1 Question 7a

When evaluating your model, we used RMSE. In the context of estimating the value of houses, what does the residual mean for an individual homeowner? How does it affect them in terms of property taxes? Discuss the cases where residual is positive and negative separately.

When the residual is positive, the model underestimates the house's value. This might mean lower property taxes for the homeowner, but they might not realize how much their house is really worth if they sell. When it's negative, the model overestimates the house's value, which could lead to higher property taxes. This is especially unfair for people with lower-cost homes since they may end up paying more than they should. Overall, when the model is wrong, it can either save homeowners money or cost them, depending on whether the predicted price is too low or too high.



---

## 0.2 Question 7b

Reflecting back on your exploration in Questions 5 and 6a, in your own words, what makes a model's predictions of property values for tax assessment purposes "fair"?

This question is open-ended and part of your answer may depend upon your specific model; we are looking for thoughtfulness and engagement with the material, not correctness.

**Hint:** Some guiding questions to reflect on as you answer the question above: What is the relationship between RMSE, accuracy, and fairness as you have defined it? Is a model with a low RMSE necessarily accurate? Is a model with a low RMSE necessarily "fair"? Is there any difference between your answers to the previous two questions? And if so, why?

A fair model should avoid large errors and shouldn't be influenced by homeowners' socio-economic status or legal biases. It's important that the residual plot is evenly and tightly distributed around 0, since unfair taxes can create serious financial troubles for many people.





### 0.3 Extra Credit Step 1: Creating Your Model

Complete the modeling steps (you can skip the cross validation step to save memory) in the cells below.

DO NOT ADD ANY EXTRA CELLS BELOW (for this part of the problem)

## 1 Please include all of your feature engineering processes

```
# for the training, validation and test sets.
# dataset_type is a flag to use as follows:
# dataset_type=1 implies the data is the training data
# dataset_type=2 implies the data is the validation data
# dataset_type = 3 implies the data is the test data

#Important Instructions:
# When processing the training data, you CAN drop any rows/data that you deem to be outliers
# When processing the validation data you CANNOT drop any rows
# When processing the test data, you CANNOT drop any rows, and you CANNOT reference "Sale Price", as it
```

In [74]: *# Modeling Step 1: Process the Data*

```
# Hint: You can either use your implementation of the One Hot Encoding Function from
#Project Part 1, or use the staff's implementation
from feature_func import *

def ohe_towncode(data):

    oh_enc = OneHotEncoder()
    oh_enc.fit(data[['Town Code']])
    dummies = pd.DataFrame(oh_enc.transform(data[['Town Code']]).todense(),
                           columns=oh_enc.get_feature_names_out(),
                           index=data.index)
    # Join the dummies to the original DataFrame
    return data.join(dummies)

def ohe_propertyclass(data):
    oh_enc = OneHotEncoder()
    oh_enc.fit(data[['Property Class']])
    dummies = pd.DataFrame(oh_enc.transform(data[['Property Class']]).todense(),
                           columns=oh_enc.get_feature_names_out(),
                           index=data.index)
    # Join the dummies to the original DataFrame
    return data.join(dummies)

## Optional: Define any helper functions you need (for example, for one-hot
# encoding, etc) above this line
```

```

def process_data_ec(df, dataset_type):

    """
    Function that includes all feature engineering processes
    for the training, validation and test sets for your extra credit model.

    Important Instructions:
    When processing the training data, you CAN drop any rows/data that you deem to be outliers
    When processing the validation data you CANNOT drop any rows (even what you deem to be an outlier)
    When processing the test data, you CANNOT drop any rows, and you CANNOT reference "Sale Price"

    dataset_type is a flag to use as follows:
        dataset_type=1 implies the data is the training data
        dataset_type=2 implies the data is the validation data
        dataset_type = 3 implies the data is the test data

    """
    data = df.copy()

    #for test dataset
    if dataset_type == 3:
        relevant_columns = []
    else:

        data = data[(data["Pure Market Filter"] == 1) & (data["Sale Price"] > 100)]

        data["Log Sale Price"] = np.log(data["Sale Price"])

        relevant_columns = ['Log Sale Price']

    # Create Log Building Square Feet column
    data["Log Building Square Feet"] = np.log(data["Building Square Feet"])

    # Create Log Estimate (Building) column
    data["Log Estimate (Building)"] = np.log(data["Estimate (Building)"] + 1)

    data["Log Age"] = np.log(data["Age"])

    data["Log Estimate (Land)"] = np.log(data["Estimate (Land)"] + 1)

    # Change Roof Material to names
    data = substitute_roof_material(data)

    # one-hot encode the roof material
    data = ohe_roof_material(data)

    data = ohe_towncode(data)

    data = ohe_propertyclass(data)

```

```

# select only relevant columns
relevant_columns += ['Log Estimate (Building)', 'Log Building Square Feet', 'Log Age', 'Log
relevant_columns += [col for col in data.columns if "Town Code" in col]
relevant_columns += [col for col in data.columns if "Property Class" in col]
# relevant_columns += [f'Property_Class_{pc}' for pc in property_classes]
data = data[relevant_columns]

return data

# Use the same original train and valid datasets from 3a (otherwise the
# validation errors aren't comparable). Don't resplit the data.

# Process the data
processed_train_ec = process_data_ec(tr, dataset_type = 1)

processed_val_ec = process_data_ec(val, dataset_type = 2)

X_train_ec = processed_train_ec.drop(columns = ['Log Sale Price'])
y_train_ec = processed_train_ec['Log Sale Price']

X_valid_ec = processed_val_ec.drop(columns = ['Log Sale Price'])
y_valid_ec = processed_val_ec['Log Sale Price']

# Take a look at the result
display(X_train_ec.head())
display(y_train_ec.head())

display(X_valid_ec.head())
display(y_valid_ec.head())

```

	Log Estimate (Building)	Log Building Square Feet	Log Age	\
21302	11.248061	6.871091	4.276666	
19451	11.510031	7.576610	4.736198	
32018	12.096431	6.891626	4.248495	
144262	12.042265	7.186901	4.007333	
197227	12.612640	7.576610	4.465908	

	Log Estimate (Land)	Roof Material_Other	Roof Material_Shake	\
21302	9.838469	0.0	0.0	
19451	10.686270	0.0	0.0	
32018	10.532123	0.0	0.0	
144262	10.270593	0.0	0.0	
197227	11.033615	0.0	0.0	

	Roof Material_Shingle/Asphalt	Roof Material_Slate	\
--	-------------------------------	---------------------	---

21302	1.0	0.0
19451	1.0	0.0
32018	1.0	0.0
144262	1.0	0.0
197227	1.0	0.0

	Roof Material_Tar&Gravel	Roof Material_Tile	...	Property Class \
21302	0.0	0.0	...	205
19451	0.0	0.0	...	205
32018	0.0	0.0	...	202
144262	0.0	0.0	...	203
197227	0.0	0.0	...	205

	Property Class_202	Property Class_203	Property Class_204 \
21302	0.0	0.0	0.0
19451	0.0	0.0	0.0
32018	1.0	0.0	0.0
144262	0.0	1.0	0.0
197227	0.0	0.0	0.0

	Property Class_205	Property Class_206	Property Class_207 \
21302	1.0	0.0	0.0
19451	1.0	0.0	0.0
32018	0.0	0.0	0.0
144262	0.0	0.0	0.0
197227	1.0	0.0	0.0

	Property Class_208	Property Class_209	Property Class_278
21302	0.0	0.0	0.0
19451	0.0	0.0	0.0
32018	0.0	0.0	0.0
144262	0.0	0.0	0.0
197227	0.0	0.0	0.0

[5 rows x 59 columns]

21302	11.738466
19451	12.100712
32018	12.577291
144262	12.502467
197227	12.449019

Name: Log Sale Price, dtype: float64

	Log Estimate (Building)	Log Building Square Feet	Log Age \
17112	11.882217	7.487174	4.043051
189337	12.110118	7.688913	4.644391
141725	11.760418	6.985642	4.127134
9776	11.569599	6.846943	4.499810
81676	11.272776	7.096721	4.158883

	Log Estimate (Land)	Roof Material_Other	Roof Material_Shake \
--	---------------------	---------------------	-----------------------

17112	10.357775	0.0	0.0
189337	11.077068	0.0	0.0
141725	10.434733	0.0	0.0
9776	9.918918	0.0	0.0
81676	10.008793	0.0	0.0

	Roof Material_Shingle/Asphalt	Roof Material_Slate \
17112	1.0	0.0
189337	1.0	0.0
141725	1.0	0.0
9776	1.0	0.0
81676	1.0	0.0

	Roof Material_Tar&Gravel	Roof Material_Tile ...	Property Class \
17112	0.0	0.0 ...	203
189337	0.0	0.0 ...	205
141725	0.0	0.0 ...	203
9776	0.0	0.0 ...	202
81676	0.0	0.0 ...	203

	Property Class_202	Property Class_203	Property Class_204 \
17112	0.0	1.0	0.0
189337	0.0	0.0	0.0
141725	0.0	1.0	0.0
9776	1.0	0.0	0.0
81676	0.0	1.0	0.0

	Property Class_205	Property Class_206	Property Class_207 \
17112	0.0	0.0	0.0
189337	1.0	0.0	0.0
141725	0.0	0.0	0.0
9776	0.0	0.0	0.0
81676	0.0	0.0	0.0

	Property Class_208	Property Class_209	Property Class_278
17112	0.0	0.0	0.0
189337	0.0	0.0	0.0
141725	0.0	0.0	0.0
9776	0.0	0.0	0.0
81676	0.0	0.0	0.0

[5 rows x 59 columns]

17112	12.100707
189337	12.460715
141725	12.165251
9776	11.982929
81676	11.719940

Name: Log Sale Price, dtype: float64

In [75]: # Run this code to make sure you haven't dropped any of the rows in the validation set

```
assert X_valid_ec.shape[0] == 33475
```

```
In [78]: # MODELING STEP 3: Evaluate the RMSE for your model
```

```
linear_model_ec = lm.LinearRegression(fit_intercept=False)

linear_model_ec.fit(X_train_ec, y_train_m3)

y_predict_train_ec = linear_model_ec.predict(X_train_ec)

y_predict_valid_ec = linear_model_ec.predict(X_valid_ec)

# Training and test errors for the model
#(in its original values before the log transform)

training_error_ec = rmse(y_predict_train_ec, y_train_ec)
validation_error_ec = rmse(y_predict_valid_ec, y_valid_ec)

training_error_ec = rmse(np.exp(y_predict_train_ec), np.exp(y_train_ec))
validation_error_ec = rmse(np.exp(y_predict_valid_ec), np.exp(y_valid_ec))

(print("Extra Credit \nTraining RMSE:$ {}\nValidation RMSE:$ {}\n"
      .format(training_error_ec, validation_error_ec))
)
```

Extra Credit

Training RMSE:\$ 165037.16802993405


Validation RMSE:\$ 164650.58983746613

```
In [79]: # Optional: Run this cell to visualize
```

```
#import plotly.graph_objects as go

fig = go.Figure([
go.Bar(x = ["Extra Credit Model"], y = [training_error_ec], name="Training RMSE"),
go.Bar(x = ["Extra Credit Model"], y = [validation_error_ec], name="Validation RMSE"),
])

fig
fig.update_yaxes(range=[140000,260000], title="RMSE")
# Feel free to update the range as needed
```



In [80]: # MODELING STEP 5: Plot 2 side-by-side residual plots for validation data

```
fig, ax = plt.subplots(1,2, figsize=(15, 5))

residuals = y_valid_ec - y_predict_valid_ec



x_plt1 = y_predict_valid_ec
y_plt1 = residuals

x_plt2 = y_valid_ec
y_plt2 = residuals

ax[0].scatter(x_plt1, y_plt1, alpha=.25)
ax[0].axhline(0, c='black', linewidth=1)
ax[0].set_xlabel(r'Predicted Log(Sale Price)')
ax[0].set_ylabel(r'Residuals: Log(Sale Price) - Predicted Log(Sale Price)');
ax[0].set_title("EC Val Data: Residuals vs. Predicted Log(Sale Price)")

ax[1].scatter(x_plt2, y_plt2, alpha=.25)
ax[1].axhline(0, c='black', linewidth=1)
ax[1].set_xlabel(r'Log(Sale Price)')
ax[1].set_ylabel(r'Residuals: Log(Sale Price) - Predicted Log(Sale Price)');
ax[1].set_title("EC Val Data: Residuals vs. Log(Sale Price)")
```

Out[80]: Text(0.5, 1.0, 'EC Val Data: Residuals vs. Log(Sale Price)')





### 1.1 Extra Credit Step 2: Explanation (Required for points on model above):

- Explain what you did to create your model. What versions did you try? What worked and what didn't?
- Comment on the RMSE and residual plots from your model. Are the residuals of your model still showing a trend that overestimates lower priced houses and underestimates higher priced houses?

**Write your answers in the text cell below**

I had played around a lot with different columns, such as Estimate(Land), Neighborhood Code, and Age Decade; however, those features didn't have a strong effect on the total RMSE. So, I just removed them and continued to add on other columns until I noticed noticeable changes. For instance, what really decreased my RMSE were property class and town codes, as once I one hot encoded and implemented them into my model, I noticed a significant decrease in the total compared to when I used the other previously mentioned features. The RMSE and residual plots definitely appear to be more significant than the ones from the previous models as well. There is still a slight trend seen in the graphs, but the dispersion is getting closer to being around 0 compared to before.

